

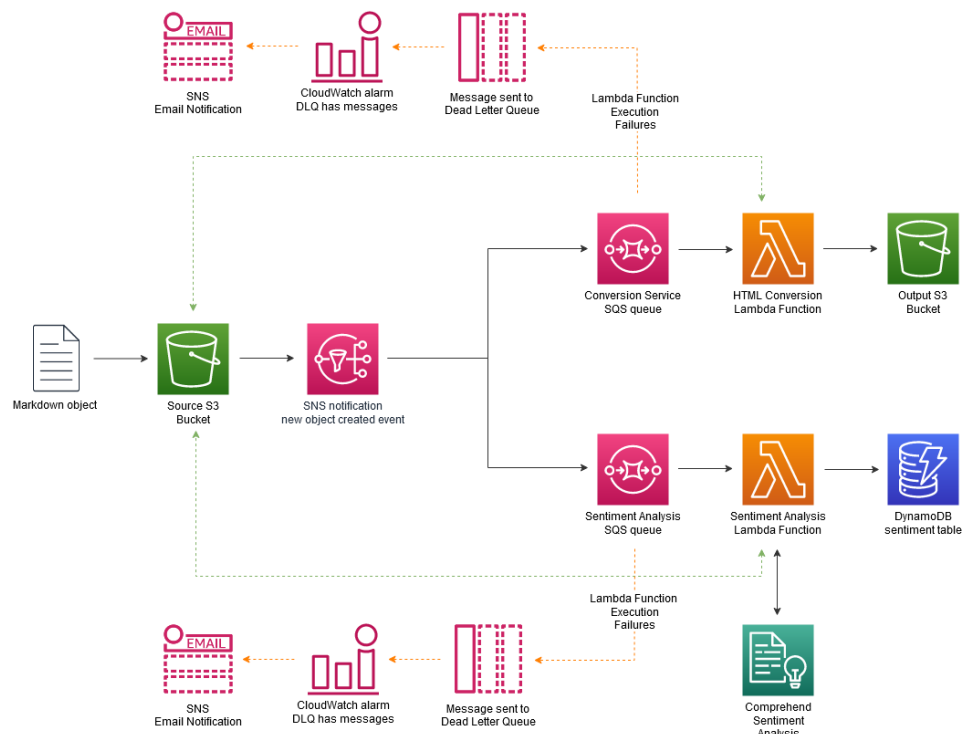
SERVERLESS IOT DATA PROCESSING

Phase 5 submission :

The Real-time File Processing reference architecture is a general-purpose, event-driven, parallel data processing architecture that uses [AWS Lambda](#). This architecture is ideal for workloads that need more than one data derivative of an object.

In this example application, we deliver notes from an interview in Markdown format to S3. S3 Events are used to trigger multiple processing flows - one to convert and persist Markdown files to HTML and another to detect and persist sentiment.

Architectural Diagram



Application Components

Event Trigger

In this architecture, individual files are processed as they arrive. To achieve this, we utilize [AWS S3 Events](#) and [Amazon Simple Notification Service](#). When an object is created in S3, an event is emitted to a SNS topic. We deliver our event to 2 separate [SQS Queues](#), representing 2 different workflows. Refer to [What is Amazon Simple Notification Service?](#) for more information about eligible targets.

Conversion Workflow

Our function will take Markdown files stored in our **InputBucket**, convert them to HTML, and store them in our **OutputBucket**. The **ConversionQueue** SQS queue captures the S3 Event JSON payload, allowing for more control of our **ConversionFunction** and better error handling. Refer to [Using AWS Lambda with Amazon SQS](#) for more details.

If our **ConversionFunction** cannot remove the messages from the **ConversionQueue**, they are sent to **ConversionDlq**, a dead-letter queue (DLQ), for inspection. A CloudWatch Alarm is configured to send notification to an email address when there are any messages in the **ConversionDlq**.

Sentiment Analysis Workflow

Our function will take Markdown files stored in our **InputBucket**, detect the overall sentiment for each file, and store the result in our **SentimentTable**.

We are using [Amazon Comprehend](#) to detect overall interview sentiment. Amazon Comprehend is a machine learning powered service that makes it easy to find insights and relationships in text. We use the Sentiment Analysis API to understand whether interview responses are positive or negative.

The Sentiment workflow uses the same SQS-to-Lambda Function pattern as the Conversion workflow.

If our **SentimentFunction** cannot remove the messages from the **SentimentQueue**, they are sent to **SentimentDlq**, a dead-letter queue (DLQ), for inspection. A CloudWatch Alarm is configured to send notification to an email address when there are any messages in the **SentimentDlq**.

Building and Deploying the Application with the AWS Serverless Application Model (AWS SAM)

This application is deployed using the [AWS Serverless Application Model \(AWS SAM\)](#). AWS SAM is an open-source framework that enables you to build serverless applications on AWS. It provides you with a template specification to define your serverless application, and a command line interface (CLI) tool.

Pre-requisites

- [AWS CLI version 2](#)
- [AWS SAM CLI \(0.41.0 or higher\)](#)
- [Docker](#)

Clone the Repository

Clone with SSH

```
git clone git@github.com:aws-samples/lambda-refarch-fileprocessing.git
```

Clone with HTTPS

```
git clone https://github.com/aws-samples/lambda-refarch-fileprocessing.git
```

Build

The AWS SAM CLI comes with abstractions for a number of Lambda runtimes to build your dependencies, and copies the source code into staging folders so that everything is ready to be packaged and deployed. The *sam build* command builds any dependencies that your application has, and copies your application source code to folders under *.aws-sam/build* to be zipped and uploaded to Lambda.

```
sam build --use-container
```

Note

Be sure to use v0.41.0 of the AWS SAM CLI or newer. Failure to use the proper version of the AWS SAM CLI will result in a `InvalidDocumentException` exception. The `EventInvokeConfig` property is not recognized in earlier versions of the AWS SAM CLI. To confirm your version of AWS SAM, run the command `sam --version`.

Deploy

For the first deployment, please run the following command and save the generated configuration file *samconfig.toml*. Please use **lambda-file-refarch** for the stack name.

```
sam deploy --guided
```

You will be prompted to enter data for *ConversionLogLevel* and *SentimentLogLevel*. The default value for each is *INFO* but you can also enter *DEBUG*. You will also be prompted for *AlarmRecipientEmailAddress*.

Subsequent deployments can use the simplified `sam deploy`. The command will use the generated configuration file *samconfig.toml*.

You will receive an email asking you to confirm subscription to the `lambda-file-refarch-AlarmTopic` SNS topic that will receive alerts should either the `ConversionDlq` SQS queue or `SentimentDlq` SQS queue receive messages.

Testing the Example

After you have created the stack using the CloudFormation template, you can manually test the system by uploading a Markdown file to the `InputBucket` that was created in the stack.

Alternatively you test it by utilising the pipeline `tests.sh` script, however the test script removes the resources it creates, so if you wish to explore the solution and see the output files and DynamoDB tables manually uploading is the better option.

Manually testing

You can use any of the `sample-xx.md` files in the repository **/tests** directory as example files. After the files have been uploaded, you can see the resulting HTML file in the output bucket of your stack. You can also view the CloudWatch logs for each of the functions in order to see the details of their execution.

You can use the following commands to copy a sample file from the provided S3 bucket into the input bucket of your stack.

```
INPUT_BUCKET=$(aws cloudformation describe-stack-resource --stack-name lambda-file-refarch --logical-resource-id InputBucket --query "StackResourceDetail.PhysicalResourceId" --output text)
aws s3 cp ./tests/sample-01.md s3://${INPUT_BUCKET}/sample-01.md
aws s3 cp ./tests/sample-02.md s3://${INPUT_BUCKET}/sample-02.md
```

Once the input files has been uploaded to the input bucket, a series of events are put into motion.

1. The input Markdown files are converted and stored in a separate S3 bucket.

```
OUTPUT_BUCKET=$(aws cloudformation describe-stack-resource --stack-name lambda-file-refarch --logical-resource-id ConversionTargetBucket --query "StackResourceDetail.PhysicalResourceId" --output text)
aws s3 ls s3://${OUTPUT_BUCKET}
```

2. The input Markdown files are analyzed and their sentiment published to a DynamoDB table.

```
DYNAMO_TABLE=$(aws cloudformation describe-stack-resource --stack-name lambda-file-refarch --logical-resource-id SentimentTable --query "StackResourceDetail.PhysicalResourceId" --output text)
aws dynamodb scan --table-name ${DYNAMO_TABLE} --query "Items[*]"
```

You can also view the CloudWatch logs generated by the Lambda functions.

Using the test script

The pipeline end to end test script can be manually executed, you will need to ensure you have adequate permissions to perform the test script actions.

- Describing stack resources
- Uploading and deleting files from the S3 input bucket
- Deleting files from the S3 output bucket
- Reading and deleting entries from the DynamoDB table

```
bash ./tests.sh lambda-file-refarch
```

While the script is executing you will see all the stages output to the command line. The samples are uploaded to the **InputBucket**, the script will then wait for files to appear in the **OutputBucket** before checking they have all been processed and the matching html file exists in the **OutputBucket**. It will also check that the sentiment for each of the files has been recorded in the **SentimentTable**. Once complete the script will remove all the files created and the entries from the **SentimentTable**.

Extra credit testing

Try uploading (or adding to ./tests if you are using the script) an oversized (> 100MB) or invalid file type to the input bucket. You can check in X-ray to explore how you can trace these kind of errors within the solution.

```
fallocate -l 110M ./tests/sample-oversize.md
```

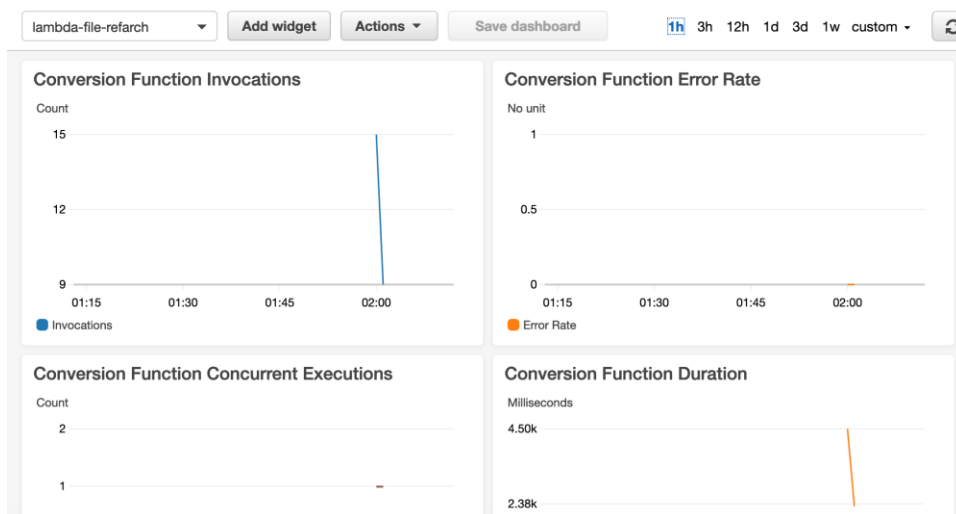
- Mac OS X command

```
mkfile 110m ./tests/sample-oversize.md
```



Viewing the CloudWatch dashboard

A dashboard is created as a part of the stack creation process. Metrics are published for the conversion and sentiment analysis processes. In addition, the alarms and alarm states are published.



Cleaning Up the Example Resources

To remove all resources created by this example, run the following command:

```
bash cleanup.sh
```

What Is Happening in the Script?

Objects are cleared out from the InputBucket and ConversionTargetBucket.

```
for bucket in InputBucket ConversionTargetBucket; do
    echo "Clearing out ${bucket}..."
    BUCKET=$(aws cloudformation describe-stack-resource --stack-name lambda-file-
refarch --logical-resource-id ${bucket} --query
"StackResourceDetail.PhysicalResourceId" --output text)
    aws s3 rm s3://${BUCKET} --recursive
    echo
done
```

The CloudFormation stack is deleted.

```
aws cloudformation delete-stack \
--stack-name lambda-file-refarch
```

The CloudWatch Logs Groups associated with the Lambda functions are deleted.

```
for log_group in $(aws logs describe-log-groups --log-group-name-prefix
'/aws/lambda/lambda-file-refarch-' --query "logGroups[*].logGroupName" --output
text); do
    echo "Removing log group ${log_group}..."
    aws logs delete-log-group --log-group-name ${log_group}
    echo
done
```

SAM Template Resources

Resources

The provided template creates the following resources:

- **InputBucket** - A S3 bucket that holds the raw Markdown files. Uploading a file to this bucket will trigger processing functions.
- **NotificationTopic** - A SNS topic that receives S3 events from the **InputBucket**.
- **NotificationTopicPolicy** - A SNS topic policy that allows the **InputBucket** to publish events to the **NotificationTopic**.

- **NotificationQueuePolicy** - A SQS queue policy that allows the **NotificationTopic** to publish events to the **ConversionQueue** and **SentimentQueue**.
- **ApplyS3NotificationLambdaFunction** - A Lambda function that adds a S3 bucket notification when objects are created in the **InputBucket**. The function is called by **ApplyInputBucketTrigger**.
- **ApplyInputBucketTrigger** - A CloudFormation Custom Resource that invokes the **ApplyS3NotificationLambdaFunction** when a CloudFormation stack is created.
- **ConversionSubscription** - A SNS subscription that allows the **ConversionQueue** to receive messages from **NotificationTopic**.
- **ConversionQueue** - A SQS queue that is used to store events for conversion from Markdown to HTML.
- **ConversionDlq** - A SQS queue that is used to capture messages that cannot be processed by the **ConversionFunction**. The *RedrivePolicy* on the **ConversionQueue** is used to manage how traffic makes it to this queue.
- **ConversionFunction** - A Lambda function that takes the input file, converts it to HTML, and stores the resulting file to **ConversionTargetBucket**.
- **ConversionTargetBucket** - A S3 bucket that stores the converted HTML.
- **SentimentSubscription** - A SNS subscription that allows the **SentimentQueue** to receive messages from **NotificationTopic**.
- **SentimentQueue** - A SQS queue that is used to store events for sentiment analysis processing.
- **SentimentDlq** - A SQS queue that is used to capture messages that cannot be processed by the **SentimentFunction**. The *RedrivePolicy* on the **SentimentQueue** is used to manage how traffic makes it to this queue.
- **SentimentFunction** - A Lambda function that takes the input file, performs sentiment analysis, and stores the output to the **SentimentTable**.
- **SentimentTable** - A DynamoDB table that stores the input file along with the sentiment.
- **AlarmTopic** - A SNS topic that has an email as a subscriber. This topic is used to receive alarms from the **ConversionDlqAlarm**, **SentimentDlqAlarm**, **ConversionQueueAlarm**, **SentimentQueueAlarm**, **ConversionFunctionErrorRateAlarm**, **SentimentFunctionErrorRateAlarm**, **ConversionFunctionThrottleRateAlarm**, and **SentimentFunctionThrottleRateAlarm**.

- **ConversionDlqAlarm** - A CloudWatch Alarm that detects when there are any messages sent to the **ConvesionDlq** within a 1 minute period and sends a notification to the **AlarmTopic**.
- **SentimentDlqAlarm** - A CloudWatch Alarm that detects when there are any messages sent to the **SentimentDlq** within a 1 minute period and sends a notification to the **AlarmTopic**.
- **ConversionQueueAlarm** - A CloudWatch Alarm that detects when there are 20 or more messages in the **ConversionQueue** within a 1 minute period and sends a notification to the **AlarmTopic**.
- **SentimentQueueAlarm** - A CloudWatch Alarm that detects when there are 20 or more messages in the **SentimentQueue** within a 1 minute period and sends a notification to the **AlarmTopic**.
- **ConversionFunctionErrorRateAlarm** - A CloudWatch Alarm that detects when there is an error rate of 5% over a 5 minute period for the **ConversionFunction** and sends a notification to the **AlarmTopic**.
- **SentimentFunctionErrorRateAlarm** - A CloudWatch Alarm that detects when there is an error rate of 5% over a 5 minute period for the **SentimentFunction** and sends a notification to the **AlarmTopic**.
- **ConversionFunctionThrottleRateAlarm** - A CloudWatch Alarm that detects when there is a throttle rate of 1% over a 5 minute period for the **ConversionFunction** and sends a notification to the **AlarmTopic**.
- **SentimentFunctionThrottleRateAlarm** - A CloudWatch Alarm that detects when there is a throttle rate of 1% over a 5 minute period for the **SentimentFunction** and sends a notification to the **AlarmTopic**.
- **ApplicationDashboard** - A CloudWatch Dashboard that displays Conversion Function Invocations, Conversion Function Error Rate, Conversion Function Throttle Rate, Conversion DLQ Length, Sentiment Function Invocations, Sentiment Function Error Rate, Sentiment Function Throttle Rate, and Sentiment DLQ Length.

License

This reference architecture sample is licensed under Apache 2.0

Python

```
import json

def process_data(event):
    """This function processes the data received from the IoT
    device."""

    # Get the sensor reading from the event data.
    sensor_reading = event["data"]["sensor_reading"]

    # Filter the data based on certain criteria.
    if sensor_reading > 100:
        # Trigger an automated routine.
        send_alert_email()

def send_alert_email():
    """This function sends an alert email."""

    # Create an email message.
    email_message = """
    Subject: Sensor alert!

    The sensor reading has exceeded the threshold.
    """

    # Send the email message.
    send_email(email_message)

# Register the function to be triggered when an event is
# received from the IoT device.
IBMCloudFunctions.register_function(process_data,
    "iot_device_event")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import datetime
```

```
%matplotlib inline
import scipy.integrate as integrate
from scipy.optimize import curve_fit
pd.options.display.max_rows = 12
```

```
data = pd.read_csv("../input/data.txt", sep = ' ', header = None, names = ['date',
'time', 'epoch', 'moteid', 'temp', 'humidity', 'light', 'voltage'])
data.head()
```

Out[2]:

	DATE	TIME	EPOCH	MOTEID	TEMP	HUMIDITY	LIGHT	VOLTAGE
0	2004-03-31	03:38:15.757551	2	1.0	122.1530	-3.91901	11.04	2.03397
1	2004-02-28	00:59:16.02785	3	1.0	19.9884	37.09330	45.08	2.69964
2	2004-02-28	01:03:16.33393	11	1.0	19.3024	38.46290	45.08	2.68742
3	2004-02-28	01:06:16.013453	17	1.0	19.1652	38.80390	45.08	2.68742
4	2004-02-28	01:06:46.778088	18	1.0	19.1750	38.83790	45.08	2.69964

```
data.shape
```

```
(2313682, 8)
```

Out[3]:

```
data.describe()
```

```
data.describe()
```

Out[4]:

	EPOCH	MOTEID	TEMP	HUMIDITY	LIGHT	VOLTAGE
count	2.313682e+06	2.313156e+06	2.312781e+06	2.312780e+06	2.219804e+06	2.313156e+06

mean	3.303993e+04	2.854412e+01	3.920700e+01	3.390814e+01	4.072110e+02	2.492552e+00
std	1.836852e+04	5.062408e+01	3.741923e+01	1.732152e+01	5.394276e+02	1.795743e-01
min	0.000000e+00	1.000000e+00	-3.840000e+01	-8.983130e+03	0.000000e+00	9.100830e-03
25%	1.757200e+04	1.700000e+01	2.040980e+01	3.187760e+01	3.956000e+01	2.385220e+00
50%	3.332700e+04	2.900000e+01	2.243840e+01	3.928030e+01	1.582400e+02	2.527320e+00
75%	4.778900e+04	4.100000e+01	2.702480e+01	4.358550e+01	5.372800e+02	2.627960e+00
max	6.553500e+04	6.540700e+04	3.855680e+02	1.375120e+02	1.847360e+03	1.856000e+01

```
data.isnull().sum()
```


Out[5]:

```
date          0
time          0
epoch         0
moteid        526
temp          901
humidity      902
light        93878
voltage       526
dtype: int64
```

In [6]:

```
data['temp'].plot()
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f83485fb1d0>



```
data.fillna(0, inplace=True)
data['epoch'].replace(regex=True, inplace=True, to_replace=r'^0-9.\-]',
value=r'')
data['epoch'] = data['epoch'].astype(int)
```

In [8]:

```
data.isnull().sum()
```

```
date      0
time      0
epoch     0
moteid    0
```

```
temp      0
humidity  0
light     0
voltage   0
dtype: int64
```

In [9]:

```
data.groupby('moteid').mean()
```

epoch	temp	humidity	light	voltage	
moteid					
0.0	22031.159696	0.000000	0.000000	0.000000	0.000000
1.0	31994.793830	35.882437	34.319280	156.575828	2.519643
2.0	36491.176404	40.201817	34.298739	212.159717	2.458436
3.0	33013.771411	33.715289	34.787174	146.049674	2.517961
4.0	32688.045167	45.464410	29.991159	155.491512	2.468170
5.0	5836.428571	0.000000	0.000000	0.000000	2.516088
...
56.0	54586.336657	20.987656	39.261422	104.099945	2.887558
57.0	1499.000000	0.000000	17.887833	0.000000	2.482020
58.0	53830.011773	21.014753	38.541575	364.111968	2.899639
6485.0	35112.000000	262.656000	0.000000	0.000000	0.393324
33117.0	30493.000000	-36.204800	0.000000	0.000000	0.144859
65407.0	18182.000000	294.251000	0.000000	0.000000	0.013056

```
data['Timestamp'] = data[['date', 'time']].apply(lambda x: '
'.join(x.astype(str)), axis=1)
new_data = data
```

In [11]:

```
data.drop(['date','time'],axis=1,inplace =True)
data.set_index(pd.to_datetime(data.Timestamp), inplace=True)
```

In [12]:

```
data[['moteid','temp','humidity','light','voltage']] =
data[['moteid','temp','humidity','light','voltage']].apply(pd.to_numeric)
```

```
data['moteid'].value_counts()
```

Out[13]:

```
31.0      65694
29.0      64391
23.0      62440
26.0      61521
22.0      60165
21.0      58525
...
0.0         526
5.0          35
57.0          3
65407.0       1
6485.0        1
33117.0        1
Name: moteid, Length: 62, dtype: int64
```

```
moteid_grp = data.groupby(['moteid'])
```

In [15]:

```
corr_id = moteid_grp.corr(method='pearson')
corr_id.fillna(0, inplace=True)
corr_id
```

epoch	humidity	light	temp	voltage		
moteid						
0.0	epoch	1.0	0.000000	0.000000	0.000000	0.000000
	humidity	0.0	0.000000	0.000000	0.000000	0.000000
	light	0.0	0.000000	0.000000	0.000000	0.000000
	temp	0.0	0.000000	0.000000	0.000000	0.000000
	voltage	0.0	0.000000	0.000000	0.000000	0.000000
1.0	epoch	1.0	-0.195395	0.014605	0.315767	-0.726957
...
33117.0	voltage	0.0	0.000000	0.000000	0.000000	0.000000
65407.0	epoch	0.0	0.000000	0.000000	0.000000	0.000000
	humidity	0.0	0.000000	0.000000	0.000000	0.000000
	light	0.0	0.000000	0.000000	0.000000	0.000000
	temp	0.0	0.000000	0.000000	0.000000	0.000000
	voltage	0.0	0.000000	0.000000	0.000000	0.000000

310 rows × 5 columns

epoch	moteid	temp	humidity	light	voltage	Timestamp	
Timestamp							
2004-03-31 03:38:15.757 551	2	1.0	122.1530	-3.91901	11.04	2.03397	2004-03-31 03:38:15.757 551
2004-02-28 00:59:16.027 850	3	1.0	19.9884	37.09330	45.08	2.69964	2004-02-28 00:59:16.027 85
2004-02-28 01:03:16.333 930	11	1.0	19.3024	38.46290	45.08	2.68742	2004-02-28 01:03:16.333 93

2004-02-28 01:06:16.013 453	17	1.0	19.1652	38.80390	45.08	2.68742	2004-02-28 01:06:16.013 453
2004-02-28 01:06:46.778 088	18	1.0	19.1750	38.83790	45.08	2.69964	2004-02-28 01:06:46.778 088

Temperature_data

de22	Node2 3	Node2 4	Node2 5	Node2 6	Node2 7	Node2 8	Node2 9					
Timest amp	Timest amp	Timest amp	Timest amp									
2004	3	1	0	18.892 649	17.888 361	18.682 962	17.909 755	17.859 044	17.673 663	18.159 542	16.981 087	17.684 646
			1	17.828 778	17.513 253	18.301 935	17.535 843	17.494 400	17.259 217	17.705 163	16.576 858	17.429 180
			2	17.478 076	17.269 043	18.038 685	17.276 184	17.222 923	16.988 568	17.443 851	16.263 400	17.193 869
			3	17.265 336	17.036 493	17.833 192	17.068 284	17.011 023	16.778 655	17.242 236	16.098 262	17.046 103
			4	16.976 067	16.728 080	17.526 390	16.730 619	16.708 245	16.424 560	16.829 928	15.745 938	16.794 936
			5	16.678 958	16.295 632	17.130 163	16.236 195	16.223 800	15.896 133	16.312 522	15.345 038	16.430 650
	
		21	18	21.580 820	20.955 944	24.125 349	32.503 942	21.184 272	21.099 364	22.016 777	20.406 817	21.141 881
			19	21.213 658	20.334 054	24.385 489	35.667 856	20.598 625	20.551 900	21.721 577	19.656 331	20.547 413
			20	20.893 717	19.951 263	24.994 137	40.132 790	20.216 815	20.256 095	21.630 513	19.215 485	20.144 796
			21	20.408 241	19.569 963	25.837 693	45.156 295	19.830 846	19.937 289	21.666 446	18.845 821	19.723 914
			22	20.171 971	19.167 160	27.153 425	51.666 562	19.390 084	19.724 252	21.884 614	18.476 750	19.241 492
			23	19.997 640	18.620 570	28.617 012	59.078 113	18.886 997	19.474 558	22.185 723	17.960 217	18.716 850

463 rows × 9 columns

In [19]:

Humidity_data

Out[19]:

				Node1 4	Node2 2	Node2 3	Node2 4	Node2 5	Node2 6	Node2 7	Node2 8	Node2 9
Timest amp	Timest amp	Timest amp	Timest amp									
2004	3	1	0	42.549 285	44.530 996	43.168 231	44.878 085	44.633 242	47.021 085	45.492 716	48.549 510	46.209 795
			1	45.355 191	45.693 293	44.364 838	46.024 384	45.775 104	48.405 524	46.959 683	49.996 614	46.909 245
			2	46.163 976	46.084 623	44.710 712	46.487 651	46.130 820	48.831 494	47.188 011	50.559 377	47.271 673
			3	46.613 439	46.469 025	45.070 823	46.854 633	46.512 194	49.220 666	47.632 574	50.843 643	47.565 288
			4	47.179 816	47.056 620	45.602 291	47.544 280	47.070 312	49.981 491	48.445 950	51.645 920	48.041 444
			5	47.897 721	48.318 876	46.838 235	49.095 143	48.524 784	51.741 095	50.494 748	52.944 230	49.184 805
		21
			18	48.074 626	47.581 688	46.317 189	48.043 332	45.830 167	48.557 737	46.804 405	48.878 659	46.799 120
			19	48.199 534	48.236 579	46.810 787	48.842 767	46.393 657	49.506 822	47.598 095	50.102 531	47.288 588
			20	48.285 621	48.535 250	47.076 702	49.253 592	46.754 700	49.895 563	48.139 502	50.566 962	47.463 996
			21	49.033 441	48.811 621	47.416 511	49.718 356	47.142 598	50.570 302	48.718 429	50.572 297	47.832 622
			22	49.483 676	49.228 436	47.773 140	50.317 610	47.654 925	51.163 117	49.398 291	50.813 785	48.433 595
			23	49.850 957	50.196 400	48.584 533	51.448 840	48.664 630	52.343 430	50.653 467	51.896 545	49.429 506

Covariance models in time for the temperature

Temperature time lags from 0 to 9

In [20]:

```
node_distance = pd.read_csv('../input/intel_nodes_distances.csv')
```

```
lists_hy = []
```

```


lists_hx = []

for z in range(10):
    # Calculation of the covariance for lag i hour for temp
    humy = [None] *45
    humx = [None] *45
    k=0
    i=0
    for first_column in Temperature_data:
        df1 = Temperature_data[first_column][:-(1+z)]
        std_test1=np.std(df1)
        j = 0
        for second_column in Temperature_data:
            if j <= i:
                df2 = Temperature_data[second_column][(1+z):]
                std_test2=np.std(df2)
                humy[k] = (np.cov(df1,df2)/(std_test1*std_test2)).item((0, 1))
                humx[k] = node_distance.iloc[i,j]
                j = j + 1
                k = k + 1
            i = i + 1
        lists_hy.append(humy)
        lists_hx.append(humx)

```

Out[21]:

<matplotlib.collections.PathCollection at 0x7f8327bc9550>



In [22]:

```


from statsmodels.tsa.stattools import acf,pacf
lag_acf = acf(data['temp'])

```

```
#Plot pACF: a
plt.subplot(121)
plt.plot(lag_acf)
```

Out[22]:

```
[<matplotlib.lines.Line2D at 0x7f8326a94470>]
```




Variation in humidity, temp, light, voltage with epoch

In [23]:

```
fig, axes = plt.subplots(nrows=1, ncols=4, figsize=(60,20))

for xcol, ax in zip(['humidity', 'temp', 'light', 'voltage'], axes):
    data.plot(kind='scatter', x='epoch', y=xcol, ax=ax, alpha=1, color='r')
```




Variation with Voltage

In [24]:

```
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(60,30))
for xcol, ax in zip(['temp', 'light', 'humidity'], axes):
```

```
data.plot(kind='scatter', x='voltage', y=xcol, ax=ax, color='b')
```



Variation with light

In [25]:




```
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(60,30))
for xcol, ax in zip(['temp', 'humidity', 'voltage'], axes):
    data.plot(kind='scatter', x='light', y=xcol, ax=ax, alpha=1, color='g')
```

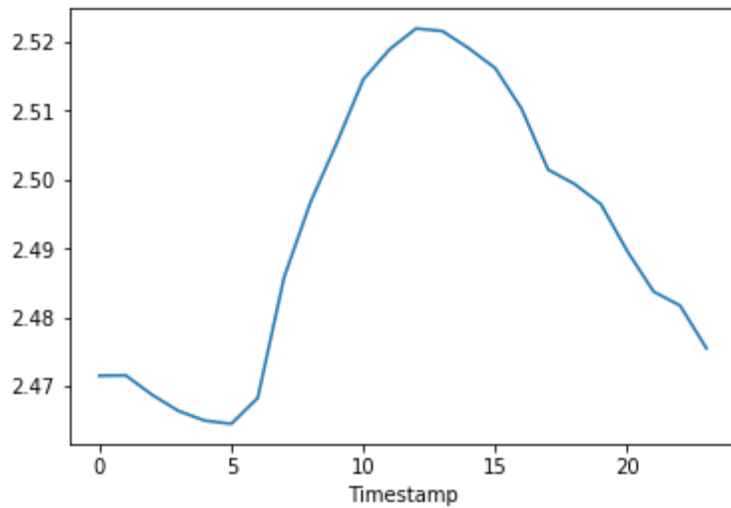
Variation with Humidity

In [26]:

```
fig, axes = plt.subplots(nrows=1, ncols=3, figsize=(60,30))
for xcol, ax in zip(['temp', 'light', 'voltage'], axes):
    data.plot(kind='scatter', x='humidity', y=xcol, ax=ax, alpha=1,
```

color='y' Variation in Tempeature, Humidity, light and Voltage over time

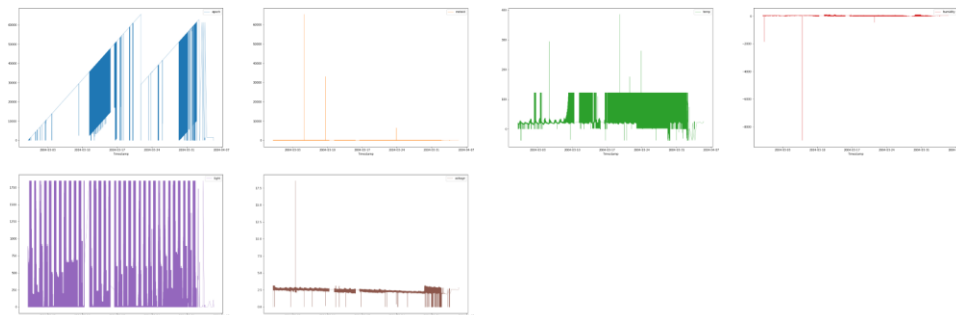




In [29]:

```
new_data.plot(subplots=True,linewidth=0.5,
              layout=(2, 4),figsize=(60, 20),
              sharex=False,
              sharey=False)
```

plt.show()



Pearson Correlation for the multivariate time series

In [30]:

```
new_data.corr(method='pearson')
```

Out[30]:

	EPOCH	MOTEID	TEMP	HUMIDITY	LIGHT	VOLTAGE
epoch	1.000000	0.006596	0.340746	-0.221278	0.041511	-0.654189

moteid	0.006596	1.000000	-0.014062	0.023541	0.029538	0.001789
temp	0.340746	-0.014062	1.000000	-0.707251	0.021948	-0.737583
humidity	-0.221278	0.023541	-0.707251	1.000000	-0.095084	0.507445
light	0.041511	0.029538	0.021948	-0.095084	1.000000	0.055835
voltage	-0.654189	0.001789	-0.737583	0.507445	0.055835	1.000000


Variation in temperature readings over time for moteid's: 21 and 22

In [31]:

```
from matplotlib import pyplot as plt
d_m21 = data.loc[data['moteid'] == 21.0]
d_m22 = data.loc[data['moteid'] == 22.0]
d_m10 = data.loc[data['moteid'] == 10.0]
fig2 = plt.figure(figsize = (15,10))
d_m21['temp'].plot(label='temperature for moteid=21.0')
d_m22['temp'].plot(label='temperature for moteid=22.0')
fig2.suptitle('Variation in temperature over time for moteid= 21.0 and 22.0',
fontsize=10)
plt.xlabel('timestamp', fontsize=10)
plt.ylabel('temperature', fontsize=10)
plt.legend()
```

Out[31]:

<matplotlib.legend.Legend at 0x7f8327b78ba8>



Anomaly Detection using moving average method

For moteid:10 and window size: 20, we calculate the mean and standard deviation of the data. If the next entry in the dataframe lies between $\text{mean}(\pm \text{sd} * 2)$, it is considered normal else it is considered an anomaly.

Anomaly can be seen by blue *

In [32]:

```
from itertools import count
import matplotlib.pyplot as plt
from numpy import linspace, loadtxt, ones, convolve
import numpy as np
import pandas as pd
import collections
from random import randint
from matplotlib import style
%matplotlib inline
def mov_average(data, window_size):

    window = np.ones(int(window_size))/float(window_size)
    return np.convolve(data, window, 'same')
def find_anomalies(y, window_size, sigma=1.0):
    avg = mov_average(y, window_size).tolist()
    residual = y - avg
    std = np.std(residual)
    return {'standard_deviation': round(std, 3),
            'anomalies_dict': collections.OrderedDict([(index, y_i) for index,
y_i, avg_i in zip(count(), y, avg)
                    if (y_i > avg_i + (sigma*std)) | (y_i < avg_i - (sigma*std))])]}
def plot_results(x, y, window_size, sigma_value=1,
                 text_xlabel="X Axis", text_ylabel="Y Axis",
                 applying_rolling_std=False):

    plt.figure(figsize=(15, 8))
    plt.plot(x, y, "k.")
    y_av = mov_average(y, window_size)
    plt.plot(x, y_av, color='green')
    plt.xlim(0, 40000)
    plt.xlabel(text_xlabel)
    plt.ylabel(text_ylabel)
    events = {}
    events = find_anomalies(y, window_size=window_size, sigma=sigma_value)
```

```

x_anom = np.fromiter(events['anomalies_dict'].keys(), dtype=int,
count=len(events['anomalies_dict']))
y_anom = np.fromiter(events['anomalies_dict'].values(),
dtype=float,count=len(events['anomalies_dict']))
plt.plot(x_anom, y_anom, "b*")
print(x_anom)
plt.grid(True)
plt.show()
x = d_m10['epoch']
Y = d_m10['temp']
plot_results(x, y=Y, window_size=50, text_xlabel="Date",
sigma_value=3,text_ylabel="temperature")

```

[23743 23751 23761 23999 24206 24302 24303 24324 24350 25415 26094
26101
26325 26336 26371 26422 26437 26549 26551 26581 26588 26622 26624
26636
26713 26723 29505 29506 29507 36012 36017 40733 40748 40775 40779
40787
40840 40841 40848 40982 41115 41201]

